

## CNN Part

```
Evaluating VGG-19...
```

```
C:\Users\nikhi\anaconda3\Lib\site-packages\keras\backend\tensorflow_backend.py:71: UserWarning: Using a deprecated name inside `keras.backend`: 'get_session' will be removed in the future, please use 'get_session_config_value' instead.
warnings.warn(msg)
C:\Users\nikhi\anaconda3\Lib\site-packages\keras\backend\tensorflow_backend.py:71: UserWarning: Using a deprecated name inside `keras.backend`: 'get_session' will be removed in the future, please use 'get_session_config_value' instead.
warnings.warn(msg)
```

```
VGG-19 Accuracy: 17.52%
```

```
Evaluating ResNet-50...
```

```
C:\Users\nikhi\anaconda3\Lib\site-packages\keras\backend\tensorflow_backend.py:71: UserWarning: Using a deprecated name inside `keras.backend`: 'get_session' will be removed in the future, please use 'get_session_config_value' instead.
warnings.warn(msg)
```

```
ResNet-50 Accuracy: 15.11%
```

```
Evaluating InceptionV4...
```

```
InceptionV4 Accuracy: 2.55%
```

```
Comparison Report:
```

```
VGG-19 Accuracy: 17.52%
```

```
ResNet-50 Accuracy: 15.11%
```

```
InceptionV4 Accuracy: 2.55%
```

## Report on Model Performance on Tiny ImageNet Dataset

### Results Overview:

- **VGG-19 Accuracy: 17.52%**
- **ResNet-50 Accuracy: 15.11%**
- **InceptionV4 Accuracy: 2.55%**

### Analysis and Comparison:

#### 1. VGG-19 Performance:

- **Architecture:** VGG-19 is a deep convolutional neural network with 19 layers and a uniform architecture design. It uses small 3x3 filters and stacks multiple convolutional layers followed by max-pooling layers.
- **Reason for Performance:**
  - The moderate accuracy of 17.52% achieved by VGG-19 indicates that the model has some ability to learn the features of the Tiny ImageNet dataset. However, the relatively low accuracy suggests challenges with

generalization due to dataset complexity, reduced data size (subset used), and the high number of classes (200 classes).

- The network's depth allows it to capture complex features, but without enough data, there may be an overfitting problem.

## **2. ResNet-50 Performance:**

- **Architecture:** ResNet-50 is a deep residual network comprising 50 layers and employs skip connections (residual blocks). These connections help in alleviating the vanishing gradient problem and enable deeper network training.
- **Reason for Performance:**
  - ResNet-50 achieved an accuracy of 15.11%, slightly lower than VGG-19. This could be due to the deeper architecture needing more data and longer training time to achieve its full potential.
  - With limited data (a subset of 2000 images) and potentially insufficient epochs, ResNet-50 might not have fully leveraged its residual connections for effective learning.

## **3. InceptionV4 Performance:**

- **Architecture:** InceptionV4 is a complex network incorporating multiple parallel convolutional operations within an Inception block to capture different feature scales and dimensions.
- **Reason for Performance:**
  - The very low accuracy (2.55%) may be attributed to challenges in adapting such a complex model to a small dataset with limited training samples.
  - InceptionV4's architecture requires substantial computational resources and data diversity to effectively generalize. Given the limited subset used for training, the model might have failed to learn meaningful patterns, resulting in poor accuracy.

## **Comparison and Conclusion:**

- **Why VGG-19 Performs Better:**
  - VGG-19's simpler architecture compared to ResNet-50 and InceptionV4 makes it more suitable for smaller datasets. Its uniform convolutional layers allow it to extract basic features without over-complicating the learning process, resulting in marginally better performance.
- **Challenges for ResNet-50 and InceptionV4:**

- Both models have deep and complex architectures designed for large-scale data. Training them on a small subset of Tiny ImageNet with 200 classes likely led to underfitting or failed convergence due to insufficient data.
- Additionally, InceptionV4's complexity makes it prone to overfitting and data sparsity issues, reflected in its poor performance.

## **RNN Part**

### **Output:**

**Running experiments for dataset size: 1000**

**Training LSTM on dataset of size 1000...**

**Epoch 1/10, Loss: 0.3432**

**Epoch 2/10, Loss: 0.3367**

**Epoch 3/10, Loss: 0.3304**

**Epoch 4/10, Loss: 0.3244**

**Epoch 5/10, Loss: 0.3185**

**Epoch 6/10, Loss: 0.3129**

**Epoch 7/10, Loss: 0.3073**

**Epoch 8/10, Loss: 0.3019**

**Epoch 9/10, Loss: 0.2964**

**Epoch 10/10, Loss: 0.2910**

**Validation Loss: 0.3131**

**Training GRU on dataset of size 1000...**

**Epoch 1/10, Loss: 0.3170**

**Epoch 2/10, Loss: 0.3040**

**Epoch 3/10, Loss: 0.2913**

**Epoch 4/10, Loss: 0.2789**

**Epoch 5/10, Loss: 0.2667**

**Epoch 6/10, Loss: 0.2548**

**Epoch 7/10, Loss: 0.2432**

**Epoch 8/10, Loss: 0.2318**

**Epoch 9/10, Loss: 0.2206**

**Epoch 10/10, Loss: 0.2096**

**Validation Loss: 0.2196**

**Training Bidirectional RNN on dataset of size 1000...**

**Epoch 1/10, Loss: 0.3312**

**Epoch 2/10, Loss: 0.3256**

**Epoch 3/10, Loss: 0.3200**

**Epoch 4/10, Loss: 0.3146**

**Epoch 5/10, Loss: 0.3092**

**Epoch 6/10, Loss: 0.3039**

**Epoch 7/10, Loss: 0.2986**

**Epoch 8/10, Loss: 0.2933**

**Epoch 9/10, Loss: 0.2880**

**Epoch 10/10, Loss: 0.2827**

**Validation Loss: 0.3056**

**Training Deep RNN on dataset of size 1000...**

**Epoch 1/10, Loss: 0.3278**

**Epoch 2/10, Loss: 0.2939**

**Epoch 3/10, Loss: 0.2625**

**Epoch 4/10, Loss: 0.2317**

**Epoch 5/10, Loss: 0.1994**

**Epoch 6/10, Loss: 0.1640**

**Epoch 7/10, Loss: 0.1248**

**Epoch 8/10, Loss: 0.0829**

**Epoch 9/10, Loss: 0.0427**

**Epoch 10/10, Loss: 0.0142**

**Validation Loss: 0.0123**

**Running experiments for dataset size: 3000**

**Training LSTM on dataset of size 3000...**

**Epoch 1/10, Loss: 0.3008**

**Epoch 2/10, Loss: 0.2947**

**Epoch 3/10, Loss: 0.2885**

**Epoch 4/10, Loss: 0.2820**

**Epoch 5/10, Loss: 0.2754**

**Epoch 6/10, Loss: 0.2685**

**Epoch 7/10, Loss: 0.2613**

**Epoch 8/10, Loss: 0.2538**

**Epoch 9/10, Loss: 0.2458**

**Epoch 10/10, Loss: 0.2374**

**Validation Loss: 0.2232**

**Training GRU on dataset of size 3000...**

**Epoch 1/10, Loss: 0.2097**

**Epoch 2/10, Loss: 0.1982**

**Epoch 3/10, Loss: 0.1868**

**Epoch 4/10, Loss: 0.1755**

**Epoch 5/10, Loss: 0.1644**

**Epoch 6/10, Loss: 0.1533**

**Epoch 7/10, Loss: 0.1424**

**Epoch 8/10, Loss: 0.1317**

**Epoch 9/10, Loss: 0.1211**

**Epoch 10/10, Loss: 0.1106**

**Validation Loss: 0.0985**

**Training Bidirectional RNN on dataset of size 3000...**

**Epoch 1/10, Loss: 0.2921**

**Epoch 2/10, Loss: 0.2861**

**Epoch 3/10, Loss: 0.2801**

**Epoch 4/10, Loss: 0.2739**

**Epoch 5/10, Loss: 0.2675**

**Epoch 6/10, Loss: 0.2608**

**Epoch 7/10, Loss: 0.2540**

**Epoch 8/10, Loss: 0.2468**

**Epoch 9/10, Loss: 0.2393**

**Epoch 10/10, Loss: 0.2314**

**Validation Loss: 0.2180**

**Training Deep RNN on dataset of size 3000...**

**Epoch 1/10, Loss: 0.0117**

**Epoch 2/10, Loss: 0.0179**

**Epoch 3/10, Loss: 0.0147**

**Epoch 4/10, Loss: 0.0085**

**Epoch 5/10, Loss: 0.0104**

**Epoch 6/10, Loss: 0.0131**

**Epoch 7/10, Loss: 0.0099**

**Epoch 8/10, Loss: 0.0079**

**Epoch 9/10, Loss: 0.0090**

**Epoch 10/10, Loss: 0.0104**

**Validation Loss: 0.0105**

**Running experiments for dataset size: 9000**

**Training LSTM on dataset of size 9000...**

**Epoch 1/10, Loss: 0.2142**

**Epoch 2/10, Loss: 0.2050**

**Epoch 3/10, Loss: 0.1951**

**Epoch 4/10, Loss: 0.1843**

**Epoch 5/10, Loss: 0.1726**

**Epoch 6/10, Loss: 0.1598**

**Epoch 7/10, Loss: 0.1457**

**Epoch 8/10, Loss: 0.1304**

**Epoch 9/10, Loss: 0.1139**

**Epoch 10/10, Loss: 0.0962**

**Validation Loss: 0.0778**

**Training GRU on dataset of size 9000...**

**Epoch 1/10, Loss: 0.0938**

**Epoch 2/10, Loss: 0.0830**

**Epoch 3/10, Loss: 0.0723**

**Epoch 4/10, Loss: 0.0616**

**Epoch 5/10, Loss: 0.0510**

**Epoch 6/10, Loss: 0.0408**

**Epoch 7/10, Loss: 0.0311**

**Epoch 8/10, Loss: 0.0223**

**Epoch 9/10, Loss: 0.0149**

**Epoch 10/10, Loss: 0.0096**

**Validation Loss: 0.0070**

**Training Bidirectional RNN on dataset of size 9000...**

**Epoch 1/10, Loss: 0.2090**

**Epoch 2/10, Loss: 0.2005**

**Epoch 3/10, Loss: 0.1913**

**Epoch 4/10, Loss: 0.1813**

**Epoch 5/10, Loss: 0.1705**

**Epoch 6/10, Loss: 0.1586**

**Epoch 7/10, Loss: 0.1456**

**Epoch 8/10, Loss: 0.1314**

**Epoch 9/10, Loss: 0.1160**

**Epoch 10/10, Loss: 0.0995**

**Validation Loss: 0.0811**

**Training Deep RNN on dataset of size 9000...**

**Epoch 1/10, Loss: 0.0092**

**Epoch 2/10, Loss: 0.0204**

**Epoch 3/10, Loss: 0.0083**

**Epoch 4/10, Loss: 0.0093**

**Epoch 5/10, Loss: 0.0134**

**Epoch 6/10, Loss: 0.0125**

**Epoch 7/10, Loss: 0.0091**

**Epoch 8/10, Loss: 0.0072**

**Epoch 9/10, Loss: 0.0086**

**Epoch 10/10, Loss: 0.0104**

**Validation Loss: 0.0095**



# Report: Performance Comparison of RNN Models on Synthetic Time Series Data

## Experiment Overview

The experiment involved training four different types of Recurrent Neural Network (RNN) models—LSTM, GRU, Bidirectional RNN, and Deep RNN—on synthetic time series data of increasing sizes. The goal was to assess the models' training and validation performance across three different dataset sizes (1000, 3000, and 9000 samples) to evaluate how well each model scales and generalizes with more data.

## Dataset Generation

- Synthetic time series data was generated using a sinusoidal function with added Gaussian noise.
- For each dataset size, the data was split into training (80%) and testing (20%) sets.
- The sequence length for the models was set to 20-time steps.

## Models Used

1. **LSTM (Long Short-Term Memory)**
2. **GRU (Gated Recurrent Unit)**
3. **Bidirectional RNN (Bidirectional LSTM)**
4. **Deep RNN (Two-layer RNN)**

## Results Summary

The results are summarized for each model across all dataset sizes, including the final training loss and validation loss after 10 epochs.

---

### 1. Dataset Size: 1000 Samples

- **LSTM:** Training Loss: 0.2910, Validation Loss: 0.3131
- **GRU:** Training Loss: 0.2096, Validation Loss: 0.2196
- **Bidirectional RNN:** Training Loss: 0.2827, Validation Loss: 0.3056
- **Deep RNN:** Training Loss: 0.0142, Validation Loss: 0.0123

### Observations:

- The Deep RNN performed exceptionally well with the smallest dataset, achieving very low training and validation losses, indicating strong overfitting to the small dataset.
  - The GRU outperformed LSTM and Bidirectional RNN in terms of achieving a lower validation loss, suggesting better generalization with the limited data.
-

## 2. Dataset Size: 3000 Samples

- **LSTM:** Training Loss: 0.2374, Validation Loss: 0.2232
- **GRU:** Training Loss: 0.1106, Validation Loss: 0.0985
- **Bidirectional RNN:** Training Loss: 0.2314, Validation Loss: 0.2180
- **Deep RNN:** Training Loss: 0.0104, Validation Loss: 0.0105

### Observations:

- All models showed improved generalization with the increased dataset size.
  - The GRU maintained its strong performance with the lowest validation loss.
  - The Deep RNN continued to exhibit very low training and validation losses, indicating potential overfitting.
- 

## 3. Dataset Size: 9000 Samples

- **LSTM:** Training Loss: 0.0962, Validation Loss: 0.0778
- **GRU:** Training Loss: 0.0096, Validation Loss: 0.0070
- **Bidirectional RNN:** Training Loss: 0.0995, Validation Loss: 0.0811
- **Deep RNN:** Training Loss: 0.0104, Validation Loss: 0.0095

### Observations:

- All models exhibited significantly lower training and validation losses with the larger dataset.
  - The GRU consistently outperformed the other models, achieving the lowest validation loss, indicating strong generalization capabilities.
  - The Deep RNN, while effective, shows signs of overfitting, as evidenced by minimal difference between training and validation loss.
- 

## Conclusion

- **Performance Comparison:** GRU consistently demonstrated the best generalization performance across all dataset sizes. This result aligns with expectations, as GRUs are known for their efficiency and ability to capture sequential dependencies effectively.
- **Effect of Increasing Dataset Size:** Increasing the dataset size improved the performance of all models, reducing validation loss and improving generalization. The larger datasets allowed models to better capture underlying patterns, reducing overfitting.

- **Model Complexity:** The Deep RNN showed strong performance with small datasets but tended to overfit, as indicated by extremely low losses. This suggests that its complexity may lead to overfitting on smaller datasets.

### **Recommendations**

- For small to medium-sized time series data, GRU is a recommended model due to its robust generalization performance.
- Bidirectional RNNs and LSTMs can also perform well but may require fine-tuning for optimal results.
- Careful monitoring of overfitting is necessary when using Deep RNNs, particularly on smaller datasets.